Reproducibility

PDF edition of preserved repository document | docs/reproducibility.md

Historical statements and test counts below retain their original scope. Consult the master report for current status.

Reproduce preserved evidence

Use the pinned foundation snapshots and WASI runtime described in docs/coding/reproducibility.md. The emergency native entry point is scripts/coding/emergency_mps.py; its original deadline is fixed and it intentionally refuses late neural reruns. Saved per-episode files checkpoint completed keys; completed CPU episodes are never deleted. Explicit historical CPU reproduction uses MINDSCAPE_DEVICE=cpu MINDSCAPE_PRECISION=float32; accelerated evaluation requires MPS and rejects silent fallback. Filesystem, tokenizer serialization, test execution, checkpoint CPU backups and Git remain CPU utilities. Neural weights, forward/generation and LoRA tensors use MPS. Batch only independent requests across episodes; decisions within each episode remain sequential.

Recompute reports with work/final-venv/bin/python scripts/coding/emergency_report.py. Restore frozen evidence into an empty directory with scripts/coding/restore_completion.py --snapshot results/final/coding_research_v2 --destination PATH; use the actual snapshot name if the mandatory locked study remained partial. The restore utility verifies every manifest hash before copying. Foundation weights are external pinned downloads, not bundled; adapters, protocol, actual datasets, source, runtime and results are bundled.

The authorized full continuation is scripts/coding/complete_research.py; it removes the superseded operational cutoff, verifies the original protocol/source/checkpoint hashes, holds a single-runner kernel lock and resumes exact missing keys. See docs/research_continuation.md. The original emergency deadline and failure records are historical.